

Monty Hall: Finite State Machines, Simulation, and Evidence

Jeremy Evert

December 2, 2025

Contents

| | |
|--|------------|
| 1 Why Games Work (and Why Monty Hall Still Haunts Our Brains) | iv |
| 1.1 Games as a learning engine | iv |
| 1.2 Why game shows became a television superpower | iv |
| 1.3 Where the money comes from (prizes, sponsors, and the network) | v |
| 1.4 Monty Hall: the man behind the doors | v |
| 1.5 <i>Let's Make a Deal</i> : a marketplace of suspense | v |
| 1.6 The modern revival | v |
| 1.7 The three-door game (and the puzzle it inspired) | v |
| 1.8 What we are building in this project | vi |
| 1.9 Where we go next | vi |
| 2 Finite State Machines (FSMs) | vii |
| 2.1 What is a finite state machine? | vii |
| 2.2 Warm-up example 1: traffic light | vii |
| 2.3 Warm-up example 2: turnstile (coin / push) | vii |
| 2.4 Warm-up example 3: login session / lockout | viii |
| 2.5 How FSMs map to code (the practical part) | viii |
| 2.6 Deterministic vs. nondeterministic (and why we mostly use deterministic in software) | ix |
| 2.7 The Monty Hall FSM (our blueprint) | x |
| 2.8 Where we go next | x |
| 3 Designing the Data-Collection FSM | xi |
| 3.1 Why a data-collection FSM at all? | xi |
| 3.2 The per-trial log (one row per game) | xi |
| 3.3 What we aggregate as we go (running totals) | xi |
| 3.4 When do we stop? | |
| Confidence, not vibes | xii |
| 3.4.1 A practical stopping rule | xiii |
| 3.4.2 Why this stopping rule matters | xiii |
| 3.4.3 A warm-up: coin tosses (confidence in action) | xiii |
| 3.4.4 The coin-toss experiment script (included in this report) | xiv |
| 3.5 Worst-case estimate of needed trials | xviii |
| 3.6 The data-collection FSM (EFSM-style) | xviii |
| 3.7 How this maps to code (preview of Chapter 4) | xix |

| | |
|---|---------------|
| 4 Designing the Code | xx |
| 4.1 Why we design the code before we write the code | xx |
| 4.2 Design goals (what we care about) | xx |
| 4.3 Suggested module layout (keep it simple) | xx |
| 4.4 The core classes (what we need and why) | xxi |
| 4.4.1 TrialRecord (one trial = one row) | xxi |
| 4.4.2 DecisionStrategy (how the player decides) | xxii |
| 4.4.3 MontyHallRules (pure logic, easy to test) | xxiii |
| 4.4.4 OutcomeAggregator (running totals + sanity checks) | xxiii |
| 4.4.5 ConfidenceMath + StoppingRule (confidence, not vibes) | xxiv |
| 4.4.6 CsvTrialLogger (simple, deterministic I/O) | xxiv |
| 4.4.7 DataCollectionFSM (the EFSM becomes a class) | xxv |
| 4.4.8 SimulationRunner (ties everything together) | xxv |
| 4.5 A testing mindset for first-year programmers | xxvi |
| 4.6 Stretch goals (optional, but excellent) | xxvi |
| 4.7 Where we go next | xxvi |
| 5 Full ChatGPT Transcript and Reflection | xxvii |
| 5.1 Transcript | xxvii |
| 5.2 What worked well | xxvii |
| 5.3 What we would do differently | xxvii |
| Bibliography | xxviii |

List of Figures

| | | |
|-----|--|------|
| 2.1 | A simple traffic light FSM. | viii |
| 2.2 | Turnstile FSM: locked/unlocked behavior. | viii |
| 2.3 | A simplified login FSM with lockout. | ix |
| 2.4 | Monty Hall FSM (Graphviz-generated). | x |
| 3.1 | Coin-toss warm-up: running estimate of \hat{p} (Heads rate) and its confidence interval tightening as n grows. | xv |
| 3.2 | Data-Collection FSM for Monty Hall simulation (Graphviz-generated). | xix |

Chapter 1

Why Games Work (and Why Monty Hall Still Haunts Our Brains)

1.1 Games as a learning engine

Games are fun for a simple reason: they turn decision-making into a tight loop. You learn the rules, take an action, observe the outcome, and adjust. That loop is fast, emotional, and memorable.

For learning programming and math, games are especially useful because they quietly force us to think in the same structures we use in computing:

- **State:** what is true right now?
- **Transitions:** what event causes the next step?
- **Rules:** what actions are allowed at each step?
- **Evidence:** how do we know our strategy is good?

This project uses a game (the Monty Hall problem) as a friendly doorway into three serious ideas: finite state machines, simulation-driven evidence, and testable software design.

1.2 Why game shows became a television superpower

Game shows fit television extremely well. A viewer can drop in mid-episode and still understand the stakes quickly. Compared to many scripted programs, game shows can also be produced efficiently, and their structure naturally leaves room for advertising.

Television history also includes an important cautionary tale: in the 1950s, several quiz shows were revealed to have been manipulated, leading to public backlash and investigations. That period helped shape later expectations around televised contests.[4]

Over time, successful daytime game shows leaned into formats that were simple, repeatable, and brand-safe. The result was programming that could attract steady audiences and therefore steady advertising dollars.

1.3 Where the money comes from (prizes, sponsors, and the network)

The financial model behind many game shows is straightforward in *structure* even when the exact numbers are opaque in *detail*:

- Networks sell advertising time against the audience the show brings in.
- Production budgets pay for staff, sets, crew, post-production, and distribution.
- Prizes may be paid directly by the show/network or supplied/offset by sponsors as part of promotional arrangements.[8]

A practical research note for students: you will often see people ask questions like “*How much did this show gross?*” or “*How much profit did the station make?*” Those figures are rarely reported cleanly at the single-show level. Networks and studios tend to report finances at the company, division, or season level. So in this chapter we focus on what we can support well: the business logic, the longevity of the format, and the way audience attention translates into revenue.

1.4 Monty Hall: the man behind the doors

Monty Hall was a Canadian-American broadcaster and producer best known as the host and co-creator of *Let’s Make a Deal*. The Television Academy notes that he appeared in more than 4,700 episodes of the show.[6]

This matters for us because it explains why the Monty Hall puzzle became culturally sticky: the show was not just a one-off novelty. It was a long-running, high-visibility format that normalized suspenseful choice as entertainment.

1.5 Let’s Make a Deal: a marketplace of suspense

Let’s Make a Deal is built around bargaining and uncertainty. Contestants (often called “traders”) make choices between known rewards and unknown possibilities hidden behind doors, curtains, or boxes. Sometimes the unknown is a great prize; sometimes it is a humorous “zonz.”

The show began in the 1960s and has had multiple runs and revivals across networks and syndication.[5] One reason the format has lasted is that it is a reliable machine for emotional moments: anticipation, risk, regret, surprise, and celebration.

Modern museum commentary emphasizes the importance of the show’s daytime audience and how that audience connects to advertising value.[7]

1.6 The modern revival

A modern version of *Let’s Make a Deal* has aired on CBS since 2009.[1] Even as sets and hosts evolve, the core ingredient remains the same: forced choices under uncertainty.

1.7 The three-door game (and the puzzle it inspired)

The classroom-friendly Monty Hall problem usually appears in a simplified, clean form:

1. There are three doors: behind one is a prize, behind the other two are goats.
2. You pick a door.
3. The host (who knows where the prize is) opens one of the other doors to reveal a goat.
4. You are offered a final choice: **stay** with your original door or **switch** to the remaining closed door.

This becomes the famous Monty Hall problem. Under the usual assumptions, switching wins with probability 2/3, while staying wins with probability 1/3.[2]

Why does this puzzle matter in a computing course? Because it is a perfect example of how humans can feel confident and still be wrong. Our intuition tends to treat the final choice as “50/50,” but the process that produced the final two doors is not symmetrical. The host’s action carries information, and the best strategy depends on modeling that action correctly.

1.8 What we are building in this project

We’ll treat the Monty Hall game as an engineered system:

- First, we specify the game precisely using a **finite state machine** (Chapter 2).
- Next, we design a **data-collection FSM** that formalizes how we will simulate and record outcomes (Chapter 3).
- Then we implement the simulation as clean **object-oriented code** with strong **unit test coverage**, and we present results with tables and plots (Chapter 4).
- Finally, we summarize what we learned (Chapter 5) and include the full ChatGPT-assisted workflow as a transcript with analysis (Chapter 6).

1.9 Where we go next

Chapter 2 introduces finite state machines and uses an FSM diagram of Monty Hall as our blueprint. Once the game is written as states and transitions, it becomes much easier to implement, test, and measure.

Chapter 2

Finite State Machines (FSMs)

2.1 What is a finite state machine?

A **finite state machine** is a simple but powerful way to model a process that moves through a limited number of situations (called **states**) based on events or inputs (called **symbols**).

At any moment:

- the machine is in **exactly one state**,
- it receives an **input event** (or a timer tick, or a button press, etc.),
- it follows a **transition** to the next state,
- and (optionally) it produces an **output/action**.

A common formal definition (for a deterministic finite automaton, DFA) is the 5-tuple $(Q, \Sigma, \delta, q_0, F)$:

- Q = set of states
- Σ = input alphabet (the set of possible events)
- δ = transition function $(Q \times \Sigma \rightarrow Q)$
- q_0 = start state
- F = set of accepting/terminal states (optional for many programming-style FSMs)

In software engineering, we often use FSMs without emphasizing “accepting states”—instead we care about: **valid transitions**, **illegal transitions**, and how to implement the state logic cleanly.

2.2 Warm-up example 1: traffic light

Before we touch Monty Hall, here is a classic FSM: a traffic light. It has a tiny state set, clear transitions, and it’s easy to test.

2.3 Warm-up example 2: turnstile (coin / push)

Another famous teaching FSM is a subway turnstile. Important lesson: **the same input can have different effects depending on the current state**. That’s the entire point of state.

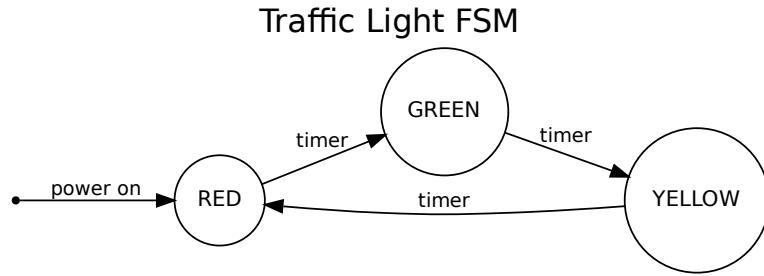


Figure 2.1: A simple traffic light FSM.

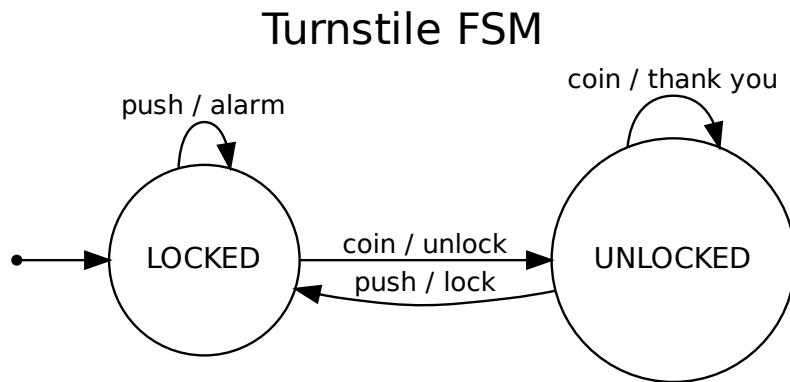


Figure 2.2: Turnstile FSM: locked/unlocked behavior.

2.4 Warm-up example 3: login session / lockout

FSMs also show up everywhere in security and application UX. A login flow is a state machine: logged out, authenticating, logged in, locked out, etc.

2.5 How FSMs map to code (the practical part)

In code, an FSM usually becomes:

- an `enum` (or strings) representing states,
- an “event” type (another enum or strings),
- a transition function: `(state, event) -> new_state`,
- optional actions performed on transitions.

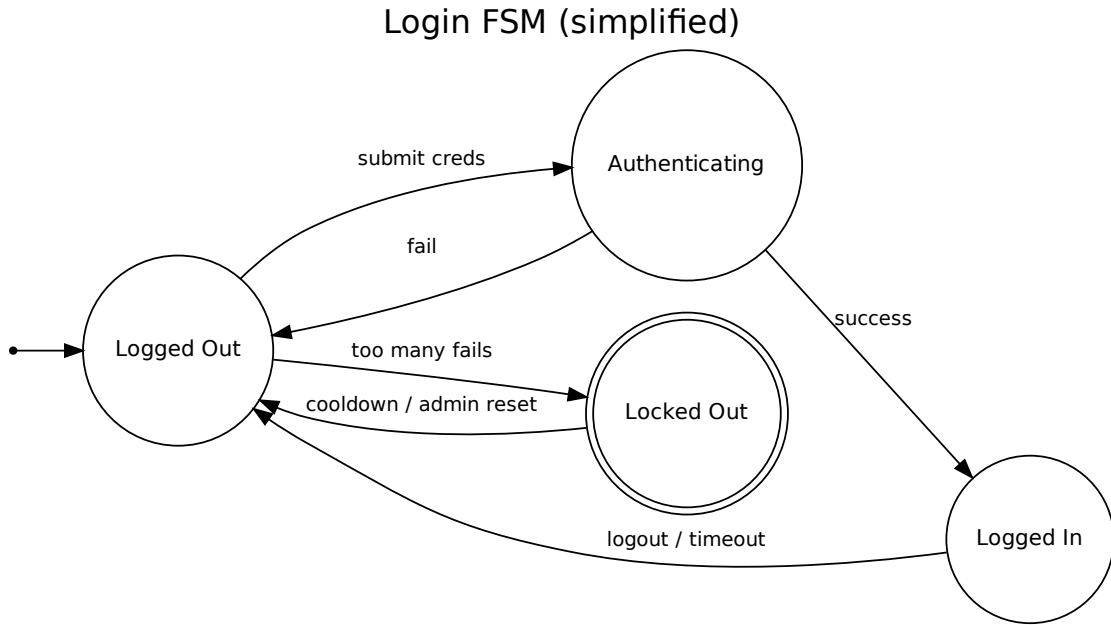


Figure 2.3: A simplified login FSM with lockout.

This approach gives you three superpowers:

1. **Clarity:** you can explain behaviour with a diagram.
2. **Testability:** you can unit-test every transition and every illegal move.
3. **Instrumentation:** you can log events and states to produce clean datasets.

2.6 Deterministic vs. nondeterministic (and why we mostly use deterministic in software)

A **deterministic** FSM has at most one outgoing transition per input symbol from each state. A **nondeterministic** FSM can have multiple possible next states for the same input. In programming projects like ours, we typically build deterministic FSMs, then model randomness as:

- probabilistic choices (e.g., random door placement),
- or hidden variables (car door, chosen door),
- or both.

When an FSM includes variables (like “car_door” or “player_choice”), many people call it an **extended finite state machine (EFSM)**: the state diagram is still the backbone, but we track a few additional values to make the model realistic.

2.7 The Monty Hall FSM (our blueprint)

Now we apply the same idea to the Monty Hall game. We will refine and expand this FSM as needed, but the key stages are: **setup** → **player picks** → **host reveals** → **player decides** → **resolve** → **reset**.




Figure 2.4: Monty Hall FSM (Graphviz-generated).

2.8 Where we go next

In Chapter 3, we design a data-collection FSM that formalizes *exactly* how simulation trials run, what gets logged, and how results flow into tables and plots. That design becomes the contract for the code and the unit tests.

Chapter 3

Designing the Data-Collection FSM (A dataset we can trust)

3.1 Why a data-collection FSM at all?

Chapter ?? argued that finite state machines give us clarity, testability, and instrumentation—especially the power to log clean datasets from a process that involves randomness. [3] Instrumenting a randomized process is the heart of empirical simulation: if we cannot exactly say when and what we recorded, we cannot trust the results.

This chapter produces two practical deliverables:

1. A precise list of **fields** that define one Monty Hall trial (one row in a CSV).
2. A clear **Data-Collection FSM** that spells out what happens, when it happens, what we log, and when we stop.

Designing this chapter first makes the FSM the *contract* for the implementation and for unit tests, ensuring both correctness and clarity.

3.2 The per-trial log (one row per game)

We record one row per simulated trial. These fields are the minimum needed to: verify randomness, debug mistakes, and compute meaningful statistics.

Each field serves a purpose:

- **trial_id** and **seed** let us exactly reproduce a particular run later.
- **prize_door**, **player_door**, **reveal_door** capture the full story of a single game.
- **decision**, **final_door**, **win**, **strategy** are the core outcomes we analyze.

3.3 What we aggregate as we go (running totals)

Along with the detailed per-trial row, we maintain running counters for three reasons:

1. Display progress without rereading the entire log.

Table 3.1: Trial log fields (one row per simulated game).

| Field | Meaning and reason to include |
|-------------|---|
| trial_id | Unique integer for traceability and reproducibility |
| seed | RNG seed (optional) to reproduce a suspicious run |
| prize_door | Where the prize was placed (should be uniform over {1,2,3}) |
| player_door | Player's initial choice (should be uniform over {1,2,3}) |
| reveal_door | Host's revealed goat door (\neq prize_door, \neq player_door) |
| decision | <code>stay</code> or <code>switch</code> (coin toss in baseline) |
| final_door | Door the player ends on after decision |
| win | 1 if final_door == prize_door, else 0 |
| strategy | <code>random</code> , <code>always_stay</code> , <code>always_switch</code> (for experiments) |

2. Build plots efficiently from incremental values.
3. Decide when we have enough trials to make trustworthy claims.

At minimum, track these four outcome buckets:

- `stay_win`
- `stay_lose`
- `switch_win`
- `switch_lose`

Add distribution checks to catch mistakes or bias:

- Counts of `prize_door` = 1,2,3
- Counts of `player_door` = 1,2,3
- Counts of `decision` stay vs switch

These sanity checks show immediately if something like door 3 never gets chosen or if randomness is skewed.

3.4 When do we stop? Confidence, not vibes

We want students to understand the difference between:

We ran a bunch of trials. vs *We ran enough trials to support a claim.*

To decide when to stop, we base the rule on a confidence interval for a proportion (the win rate). Let \hat{p} be the observed win rate for a strategy after n trials. A widely used approximate margin of error (often called the halfwidth) is:

$$\text{halfwidth} = z \sqrt{\frac{\hat{p}(1 - \hat{p})}{n}}$$

Here:

- \hat{p} is the observed win rate.
- n is the number of trials for that strategy.
- z is the z-score that matches the chosen confidence level:
 - about 1.645 for 90% confidence,
 - about 1.96 for 95%,
 - about 2.576 for 99%.

3.4.1 A practical stopping rule

Choose two things up front:

- A confidence level: 90%, 95%, or 99%.
- A tolerance ϵ : how large a halfwidth you are willing to accept. Example: $\epsilon = 0.01$ means you want the halfwidth to be at most $\pm 1\%$ around the observed win rate.

Stop the simulation only when **both** strategies' win-rate halfwidths are no larger than ϵ :

- halfwidth of switch win rate $\leq \epsilon$
- halfwidth of stay win rate $\leq \epsilon$

This makes the simulation self-aware: it continues until the estimates are sharp enough to be worth graphing or reporting.

3.4.2 Why this stopping rule matters

- If you stop too early, your estimates are noisy; the interval is wide.
- If you require very high confidence or very small tolerance, you may need many trials.
- The rule explicitly ties the stopping decision to the precision of the estimate, not to an arbitrary trial count.

3.4.3 A warm-up: coin tosses (confidence in action)

A coin toss is the simplest possible version of what we are doing in Monty Hall. Each toss is a Bernoulli trial (success/failure). If we define:

- **success** = Heads,
- **failure** = Tails,
- and \hat{p} = the observed fraction of Heads,

then the true probability is $p = 0.5$, but our estimate \hat{p} starts off wildly unstable. After one toss, \hat{p} is either 0 or 1. After two tosses, \hat{p} can be 0, 0.5, or 1. As n grows, \hat{p} tends to drift toward 0.5 (law of large numbers), but the key idea is:

We do not just want \hat{p} to be near 0.5; we want to know how confident we are in that estimate.

The same halfwidth formula applies:

$$\text{halfwidth} = z \sqrt{\frac{\hat{p}(1 - \hat{p})}{n}}$$

Two big takeaways fall right out of this equation:

1. **More trials shrink uncertainty.** Halfwidth scales like $1/\sqrt{n}$, so gains are fast early and slower later.
2. **Worst-case uncertainty happens near $\hat{p} = 0.5$.** Because $\hat{p}(1 - \hat{p})$ is maximized at 0.25, coin flips are a great “stress test” for sample size.

In the worst case ($\hat{p} \approx 0.5$), the halfwidth is approximately:

$$\text{halfwidth} \approx z \sqrt{\frac{0.25}{n}} = \frac{z}{2\sqrt{n}}$$

So if we want a 95% confidence interval with tolerance $\epsilon = 0.01$ (about $\pm 1\%$), we solve:

$$\frac{z}{2\sqrt{n}} \leq \epsilon \quad \Rightarrow \quad n \geq \frac{z^2}{4\epsilon^2}$$

For 95% confidence, $z \approx 1.96$, so:

$$n \gtrsim \frac{(1.96)^2}{4(0.01)^2} \approx 9604.$$

That number is the whole lesson: if we only flip a coin 50 times and get 60% Heads, that does *not* mean the coin is biased—it usually means the sample is still small and the confidence interval is still wide.

Figure 3.1 shows (1) the running estimate \hat{p} approaching 0.5 and (2) how the confidence interval “whiskers” tighten as n increases.

3.4.4 The coin-toss experiment script (included in this report)

To regenerate Figure 3.1, run:

```
python scripts/coin_toss_ci_demo.py
```

The script writes `figures/coin_toss_ci_demo.pdf`.

File: `scripts/coin_toss_ci_demo.py`

```
#!/usr/bin/env python3
"""
coin_toss_ci_demo.py
```

Generate a simple, classroom-friendly visualization of:

- 1) the running estimate $p_hat = (\#Heads)/n$
- 2) the confidence interval “whiskers” that tighten as n grows

Output:

```
figures/coin_toss_ci_demo.pdf
```




Figure 3.1: Coin-toss warm-up: running estimate of \hat{p} (Heads rate) and its confidence interval tightening as n grows.

This is intentionally small and dependency-light:

- Python 3.x
- matplotlib

Usage:

```
python scripts/coin_toss_ci_demo.py
python scripts/coin_toss_ci_demo.py --n 2000 --confidence 0.95 --seed 123
"""
```

```
from __future__ import annotations

import argparse
import math
from pathlib import Path
from typing import List, Tuple

import random
```

```

import matplotlib.pyplot as plt

def z_value(confidence: float) -> float:
    """
    Two-sided z critical values (common classroom defaults).
    If you want arbitrary confidence->z, you'd typically use scipy.stats,
    but we keep this lightweight and explicit.
    """
    if abs(confidence - 0.90) < 1e-9:
        return 1.645
    if abs(confidence - 0.95) < 1e-9:
        return 1.96
    if abs(confidence - 0.99) < 1e-9:
        return 2.576
    raise ValueError("confidence must be one of: 0.90, 0.95, 0.99")

def halfwidth(p_hat: float, n: int, z: float) -> float:
    """
    Approximate halfwidth of a CI for a proportion:
    z * sqrt(p_hat * (1 - p_hat) / n)

    Guarded for n=0.
    """
    if n <= 0:
        return float("inf")
    return z * math.sqrt((p_hat * (1.0 - p_hat)) / n)

def run_coin_tosses(n: int, rng: random.Random) -> Tuple[List[int], List[float], List[float], List[float]]:
    """
    Simulate n fair coin tosses.
    Returns:
        ns: [1..n]
        p_hats: running estimate of P(Heads)
        lows: p_hat - halfwidth
        highs: p_hat + halfwidth
    """
    heads = 0
    ns: List[int] = []
    p_hats: List[float] = []
    lows: List[float] = []
    highs: List[float] = []
    return ns, p_hats, lows, highs

def main() -> None:

```

```

parser = argparse.ArgumentParser(description="Coin toss CI warm-up plot generator.")
parser.add_argument("--n", type=int, default=2000, help="Number of tosses to simulate.")
parser.add_argument("--confidence", type=float, default=0.95, choices=[0.90, 0.95, 0.99],
                    help="Confidence level (two-sided): 0.90, 0.95, or 0.99.")
parser.add_argument("--seed", type=int, default=None, help="Optional RNG seed for reproduction")
parser.add_argument("--out", type=str, default="figures/coin_toss_ci_demo.pdf",
                    help="Output PDF path.")

args = parser.parse_args()

rng = random.Random(args.seed)
z = z_value(args.confidence)

# Running stats
heads = 0
ns: List[int] = []
p_hats: List[float] = []
lows: List[float] = []
highs: List[float] = []

for i in range(1, args.n + 1):
    toss_is_heads = (rng.random() < 0.5)
    if toss_is_heads:
        heads += 1

    p_hat = heads / i
    hw = halfwidth(p_hat, i, z)

    ns.append(i)
    p_hats.append(p_hat)
    lows.append(max(0.0, p_hat - hw))
    highs.append(min(1.0, p_hat + hw))

# Worst-case sample size estimate (p_hat ~ 0.5)
# halfwidth ~= z / (2*sqrt(n)) <= epsilon => n >= z^2 / (4*epsilon^2)
# We print this as a nice "sense of scale" note for students.
epsilon = 0.01
n_needed = (z * z) / (4.0 * epsilon * epsilon)

out_path = Path(args.out)
out_path.parent.mkdir(parents=True, exist_ok=True)

plt.figure()
plt.plot(ns, p_hats, label=r"\hat{p} (Heads rate)")
plt.plot(ns, lows, label="CI lower")
plt.plot(ns, highs, label="CI upper")
plt.axhline(0.5, linestyle="--", label="True p = 0.5")
plt.title(f"Coin Toss CI Tightening (confidence={args.confidence:.2f}, n={args.n})")
plt.xlabel("Number of tosses (n)")

```

```

plt.ylabel("Probability")
plt.ylim(0.0, 1.0)
plt.legend()

# Add a small annotation about worst-case n for +/-1%
plt.text(
    x=max(1, args.n // 10),
    y=0.08,
    s=f"Worst-case n for +/-1% at this confidence ~={n_needed:.0f}",
)

plt.tight_layout()
plt.savefig(out_path, format="pdf")
print(f"Wrote: {out_path}")
print(f"Worst-case n for +/-1% at confidence={args.confidence:.2f}: {n_needed:.0f}")

if __name__ == "__main__":
    main()

```

3.5 Worst-case estimate of needed trials

Because $\hat{p}(1 - \hat{p})$ is largest when $\hat{p} = 0.5$, we can derive a conservative bound on how many trials might be necessary:

$$n \geq \frac{z^2}{4\epsilon^2}$$

Derivation sketch:

- Replace $\hat{p}(1 - \hat{p})$ by its maximum value $1/4$.
- Solve the inequality $z\sqrt{\frac{1/4}{n}} \leq \epsilon$ for n .

Example: for 99% confidence ($z \approx 2.576$) and $\epsilon = 0.01$,

$$n \gtrsim \frac{(2.576)^2}{4(0.01)^2} \approx 16588.$$

This is intentionally large to demonstrate why a few hundred trials can still be noisy for tight tolerances. It is a worst-case estimate; actual needed n could be smaller if the observed \hat{p} is far from 0.5, but it gives students a sense of scale.

3.6 The data-collection FSM (EFSM-style)

When we track extra variables such as `prize_door` and `player_door`, we're really building an *extended* FSM, or EFSM. Our Data-Collection FSM uses the same clean state backbone as in Chapter ??, but now annotates each step with:

- Which variables are set at that step.

- What gets logged.
- What gets aggregated.
- When the stopping rule is evaluated.




Figure 3.2: Data-Collection FSM for Monty Hall simulation (Graphviz-generated).

3.7 How this maps to code (preview of Chapter 4)

The FSM-to-code recipe from Chapter ?? still applies:

- State enumeration,
- Events,
- Transition function,
- Actions on transitions.

The new difference is that actions include:

- Writing a trial row to CSV,
- Updating running counters,
- Computing halfwidths and checking whether to stop.

Designing this chapter first makes the FSM the contract for the implementation and for unit tests, ensuring both correctness and clarity.

Chapter 4

Designing the Code (Classes, responsibilities, and tests)

4.1 Why we design the code before we write the code

Chapter 3 defined our contract:

- one **trial row** has specific fields (Table 3.1),
- we update **running totals** as we go,
- and we **stop** only when our confidence-interval halfwidths are small enough.

That contract makes implementation much easier, because now every class we create has a job that can be tested.

This chapter proposes an object-oriented design that is intentionally *heavier* than strictly necessary, because the goal for first-year programmers is practice: **creating classes, creating objects, and testing them**.

4.2 Design goals (what we care about)

Our design has four priorities:

1. **Correctness:** the simulation obeys the rules (legal reveal door, correct final door, correct win/lose).
2. **Traceability:** every simulated trial becomes one row with the required fields. (Table 3.1)
3. **Testability:** each piece of logic is small enough to unit test in isolation.
4. **Reproducibility:** runs can be repeated using the same RNG seed.

4.3 Suggested module layout (keep it simple)

You can organize your code in a single folder without getting fancy:

```

monty_hall/
    __init__.py
    models.py          # TrialRecord, small enums
    rng.py             # RNG wrapper (optional)
    strategies.py     # DecisionStrategy + implementations
    rules.py           # MontyHallRules (pure functions)
    fsm.py             # DataCollectionFSM + states
    stats.py           # OutcomeAggregator + confidence math
    io_csv.py          # CsvTrialLogger
    runner.py          # SimulationRunner (ties it all together)

tests/
    test_rules.py
    test_strategies.py
    test_stats.py
    test_fsm.py
    test_io_csv.py

```

If your course setup prefers fewer files, you can merge some of these, but keep the class boundaries.

4.4 The core classes (what we need and why)

Chapter 3 tells us what a trial must record (trial_id, seed, prize_door, player_door, reveal_door, decision, final_door, win, strategy). :contentReference[oaicite:2]index=2 It also tells us we must track running outcome buckets and evaluate a stopping rule based on CI halfwidth. :contentReference[oaicite:3]index=3

Table 4.1 lists the recommended classes and their responsibilities.

Table 4.1: Recommended classes and responsibilities.

| Class | Primary responsibility |
|-------------------|---|
| TrialRecord | One row of trial data (exactly the Chapter 3 fields) |
| DecisionStrategy | Choose <code>stay</code> or <code>switch</code> (pluggable policies) |
| MontyHallRules | Pure rule logic: place prize, legal reveal door, finalize door, win/lose |
| OutcomeAggregator | Running totals: stay/switch win/lose + distribution checks |
| ConfidenceMath | z-values + CI halfwidth calculation |
| StoppingRule | Decide if we can stop (<code>halfwidth(stay)</code> and <code>halfwidth(switch)</code> \leq epsilon) |
| CsvTrialLogger | Write header + append <code>TrialRecords</code> to CSV |
| DataCollectionFSM | Execute the EFSM steps in the correct order and produce a <code>TrialRecord</code> |
| SimulationRunner | Loop until stop rule, update stats, write logs, return summary |

4.4.1 TrialRecord (one trial = one row)

Purpose: represent the exact trial fields from Chapter 3. :contentReference[oaicite:4]index=4

Suggested fields:

- `trial_id`: int
- `seed`: Optional[int]

- `prize_door`: int
- `player_door`: int
- `reveal_door`: int
- `decision`: str (“stay” or “switch”)
- `final_door`: int
- `win`: int (1 or 0)
- `strategy`: str

Suggested methods:

- `validate()` -> None: raises `ValueError` if fields are out of range or inconsistent.
- `to_csv_row()` -> `List[str]`: stable column order for CSV writing.
- `csv_header()` -> `List[str]`: a single canonical header list.

Unit tests:

- `test_trialrecord_validate_good()`: a known-valid record passes.
- `test_trialrecord_validate_bad_doors()`: door not in {1,2,3} fails.
- `test_trialrecord_validate_decision()`: decision must be “stay” or “switch”.
- `test_trialrecord_to_csv_row_order()`: row matches header order exactly.

4.4.2 DecisionStrategy (how the player decides)

Purpose: separate *policy* from *game logic*. A strategy chooses stay/switch, but does not touch doors.

Suggested base class:

- `choose()` -> str
- `name()` -> str

Concrete strategies:

- `AlwaysStay`: always returns “stay”
- `AlwaysSwitch`: always returns “switch”
- `CoinFlipStrategy`: uses RNG to choose (baseline described in Chapter 3) :contentReference[oaicite:5]index=5

Unit tests:

- `test_always_stay()`: 100 calls, always “stay”.
- `test_always_switch()`: 100 calls, always “switch”.
- `test_coinflip_deterministic_seed()`: with a fixed seed, first N decisions match expected sequence.
- `test_coinflip_only_two_outputs()`: outputs are only stay/switch (no surprises).

4.4.3 MontyHallRules (pure logic, easy to test)

Purpose: rules should be *pure functions* whenever possible: given inputs, return outputs; no I/O.

Suggested methods:

- `place_prize(rng) -> int`: returns 1..3
- `player_pick(rng) -> int`: returns 1..3
- `legal_reveal_door(prize_door, player_door, rng) -> int`: must not equal prize_door or player_door. :contentReference[oaicite:6]index=6
- `final_door(player_door, reveal_door, decision) -> int`: stay keeps player_door; switch chooses the remaining closed door. :contentReference[oaicite:7]index=7
- `did_win(final_door, prize_door) -> int`: returns 1 or 0. :contentReference[oaicite:8]index=8

Unit tests (these are the big ones):

- `test_legal_reveal_door_never_prize_or_player()`: for all 9 (prize, player) combos, reveal is legal.
- `test_legal_reveal_door_unique_when_player_picked_prize()`: if prize==player, reveal has 2 valid goats; ensure reveal is one of them.
- `test_final_door_stay()`: final == player_door.
- `test_final_door_switch()`: final is the only door not equal to player_door or reveal_door.
- `test_did_win_truth_table()`: verify win=1 iff final==prize.

4.4.4 OutcomeAggregator (running totals + sanity checks)

Purpose: maintain the running counters identified in Chapter 3: `stay_win`, `stay_lose`, `switch_win`, `switch_lose`, plus distributions. :contentReference[oaicite:9]index=9

Suggested fields:

- outcome buckets: `stay_win`, `stay_lose`, `switch_win`, `switch_lose`
- distributions: counts for prize_door 1/2/3, player_door 1/2/3, decision stay/switch

Suggested methods:

- `update(record: TrialRecord) -> None`
- `n_stay() -> int`, `n_switch() -> int`
- `p_stay() -> float`, `p_switch() -> float` (win rates)

Unit tests:

- `test_update_bucket_stay_win()`: one record increments only the correct bucket.
- `test_update_bucket_switch_lose()`: similarly for another bucket.
- `test_distributions_increment()`: prize/player/decision distributions increment correctly.
- `test_p_stay_and_p_switch_divide_by_zero_safe()`: 0 trials returns 0.0 or raises a clean error (choose one policy).

4.4.5 ConfidenceMath + StoppingRule (confidence, not vibes)

Chapter 3 defines halfwidth and uses it to stop when both strategies are precise enough. :contentReference[oaicite:10]index=10

ConfidenceMath methods:

- `z_value(confidence: float) -> float (90/95/99)`
- `halfwidth(p_hat: float, n: int, z: float) -> float` :contentReference[oaicite:11]index=11

StoppingRule fields:

- `confidence: float`
- `epsilon: float`

StoppingRule method:

- `should_stop(stats: OutcomeAggregator) -> bool`
- This implements: `halfwidth(stay) <= epsilon AND halfwidth(switch) <= epsilon.` :contentReference[oaicite:12]index=12

Unit tests:

- `test_z_value_known_constants()`: 0.95 maps to about 1.96 (etc.).
- `test_halfwidth_n_zero()`: returns inf (or a clear policy) :contentReference[oaicite:13]index=13
- `test_should_stop_false_early()`: small n should not stop.
- `test_should_stop_true_synthetic()`: create a fake aggregator with large n and p near 0.5 so halfwidth is tiny.

4.4.6 CsvTrialLogger (simple, deterministic I/O)

Purpose: write exactly one header row and then one row per trial record.

Suggested methods:

- `__init__(path: str)`
- `write_header() -> None`
- `append(record: TrialRecord) -> None`

Unit tests:

- `test_writes_header_once()`: call twice; file still has one header row.
- `test_append_writes_one_line()`: line count increments by one.
- `test_append_column_count()`: each row has same number of columns as header.

4.4.7 DataCollectionFSM (the EFSM becomes a class)

Chapter 3 lists the EFSM steps (place prize, pick door, reveal, decide, finalize, resolve, log, update, stop). :contentReference[oaicite:14]index=14 We can implement that flow as a class so the diagram stays the boss.

Suggested fields (context variables):

- trial_id: int
- seed: Optional[int]
- prize_door: Optional[int]
- player_door: Optional[int]
- reveal_door: Optional[int]
- decision: Optional[str]
- final_door: Optional[int]
- win: Optional[int]
- strategy_name: str

Suggested methods (beginner-friendly version):

- run_one_trial(...) -> TrialRecord: executes the EFSM stages in order and returns a record.
- If you want *extra* FSM practice: step(event) -> None and an internal state enum.

Unit tests:

- test_run_one_trial_produces_valid_record(): record.validate() passes.
- test_run_one_trial_reveal_legal(): reveal != prize and reveal != player.
- test_run_one_trial_final_door_legal(): final in {1,2,3} and not equal to reveal when switching.
- test_seed_reproducibility(): with fixed seed, the full record matches expected.

4.4.8 SimulationRunner (ties everything together)

Purpose: this is the orchestration layer:

1. run a trial,
2. log it,
3. update stats,
4. check stopping rule,
5. repeat until done.

This matches the Chapter 3 “Write Trial Row - i , Update Counters - i , Update Stopping Stats - i Stopping Rule” loop. :contentReference[oaicite:15]index=15

Suggested methods:

- `run()` -> `dict`: returns a summary: totals, win rates, and stop reason.

Unit tests:

- `test_runner_stops_eventually()`: with a relaxed epsilon (e.g., 0.20), it stops quickly.
- `test_runner_writes_n_rows()`: if you run exactly N fixed trials (optional mode), CSV has N rows.
- `test_runner_stats_match_rows()`: aggregator totals equal number of logged records.

4.5 A testing mindset for first-year programmers

Every unit test should follow the same three-beat rhythm:

1. **Arrange:** set up inputs (especially a fixed RNG seed).
2. **Act:** call exactly one method.
3. **Assert:** verify one clear outcome.

Two friendly rules:

- If a method is hard to test, it is probably doing too much.
- Put randomness behind an injected RNG so tests can be deterministic.

4.6 Stretch goals (optional, but excellent)

If a team finishes early, here are upgrades that deepen learning without changing the core design:

- Add a `FixedTrialSequenceRNG` for tests that need predictable outputs.
- Add more strategies (e.g., `SwitchUnlessPickedDoor1`).
- Produce a summary CSV with final win rates and CI halfwidths.
- Add a “max trials” safety cap so experiments never run forever.

4.7 Where we go next

With this design in place, the implementation becomes mostly mechanical: each class has a small job, each method has unit tests, and the full program is just these parts working together. We are now ready to implement the code, generate CSV datasets, and create plots that support a real empirical claim.

Chapter 5

Full ChatGPT Transcript and Reflection

5.1 Transcript

5.2 What worked well

Bullet points: prompt clarity, iteration loop, testing-as-design, diagram-first thinking, etc.

5.3 What we would do differently

How you'd tighten prompts, add constraints, verify claims, and improve reproducibility.

Bibliography

- [1] CBS. Let's make a deal. https://www.cbs.com/shows/lets_make_a_deal/. Accessed: 2025-12-01.
- [2] Encyclopaedia Britannica. Monty hall problem. <https://www.britannica.com/topic/Monty-Hall-problem>. Accessed: 2025-12-01.
- [3] NIST. Fsm notes in this project: clarity, testability, instrumentation. See Chapter 2 of this report. Internal project reference.
- [4] PBS American Experience. The aftermath of the quiz show scandal. <https://www.pbs.org/wgbh/americanexperience/features/quizshow-aftermath-quiz-show-scandal/>. Accessed: 2025-12-01.
- [5] Television Academy. Let's make a deal. <https://interviews.televisionacademy.com/shows/lets-make-a-deal>. Accessed: 2025-12-01.
- [6] Television Academy. Monty hall. <https://www.televisionacademy.com/bios/monty-hall>. Accessed: 2025-12-01.
- [7] The Strong National Museum of Play. Let's make a deal still a big deal. <https://www.museumofplay.org/blog/lets-make-a-deal-still-a-big-deal/>. Accessed: 2025-12-01.
- [8] Wrapbook. How to account for prize money in your unscripted game show. <https://www.wrapbook.com/blog/unscripted-game-show-prize-money>. Accessed: 2025-12-01.